

Time and logical clocks

Paulo Sérgio Almeida



Universidade do Minho

Physical time

Logical time

Physical time

- ▶ Time reported by our physical clocks
 - ▶ mechanical clock, digital clock, atomic clock, GPS
- ▶ One could be tempted to use physical time to make decisions
 - ▶ compare events
 - ▶ compare file versions
 - ▶ decide who gets a lock
- ▶ But clocks cannot be fully synchronized
 - ▶ comparing clocks from different nodes is dangerous
 - ▶ even a small error may have serious consequences

It is worse: there is not really a global time


There is no global time

- ▶ From special relativity
 - ▶ time in different referentials appears to pass differently
 - ▶ each observer sees time at others passing slower
- ▶ From general relativity
 - ▶ time flows at different speeds according to gravitation
 - ▶ time passes slower near a dense object
- ▶ We could compute some *hyperplane of simultaneity*
 - ▶ to obtain comparable times at different points in space
 - ▶ like GPS does – but it is amazingly complex
 - ▶ even if very good approximation ...
 - ▶ ...we should not depend on it for correctness

Play with spacetime diagrams

<https://www.geogebra.org/m/HYD7hB9v>


Past and future light cones



- ▶ Only events in the past light cone of o , like a , can influence o
- ▶ Only events in the future light cone of o , like d , can be influenced by o
- ▶ Events, like b , are *spacelike* relative to o
 - ▶ cannot influence or be influenced by o
 - ▶ comparing their times, as t_b and t_o , is meaningless
- ▶ Only comparing times of a *timelike* interval makes sense
 - ▶ i.e., compare an event with others from its past or future
 - ▶ $t_a < t_o$ and $t_o < t_d$ for all observers (reference frames)

Potential causality in distributed systems

- ▶ Light cones define the boundary of potential causality
- ▶ Distributed systems are limited to message passing
- ▶ Actual message passing defines potential causality
- ▶ Consider a single message from node i to j



- ▶ Events c and f cannot influence each other; same for d and f
- ▶ But $t_c < t_f$ and $t_d > t_f$ (assuming some reference frame)
- ▶ Event a can influence f , but like for c , $t_a < t_f$
- ▶ How can we distinguish the different roles of a and c with regards to f ?

Physical time is not worth (much)

- ▶ There is really no global physical time
- ▶ Achieving approximate global “background of time” is costly
- ▶ Comparing physical times is not enough anyway
- ▶ We need something else, to know how causality propagates

Physical time


Logical time

Basic logical time properties

- ▶ Logical time is based on data updated by programs
 - ▶ such as a counter being incremented
 - ▶ without depending on any physical device
- ▶ It can use
 - ▶ local updates
 - ▶ propagation in messages
- ▶ It should mimic some properties of physical time
 - ▶ events from the past light cone should have smaller times
 - ▶ events from the future light cone should have greater times
 - ▶ comparison to other events can be meaningless

Happens-before

- ▶ Lamport defined a *happens-before* relation (\rightarrow), given by:
 1. an event is related to future events from the same process
 2. a send happens-before the respective receive
 3. if $a \rightarrow b$ and $b \rightarrow c$, then $a < c$ (transitive closure)
- ▶ $e_1 \rightarrow e_2$ when there is a path of communication from e_1 to e_2
- ▶ In this example $a \rightarrow h$ because $a \rightarrow b$, $b \rightarrow d$, $d \rightarrow e$, $e \rightarrow g$, $g \rightarrow h$.



Two events related by happens-before define a timelike interval

Concurrent events

- ▶ Happens-before is a strict partial order
- ▶ Incomparable (unrelated) events are said to be concurrent

$$a \parallel b \Leftrightarrow a \not\rightarrow b \wedge b \not\rightarrow a$$

- ▶ Two concurrent events can define a spacelike interval
 - ▶ being physically impossible to influence each other
- ▶ Or they can define a timelike interval (but we don't know)
 - ▶ and they could influence each other by a *hidden channel*
- ▶ In a closed world assumption, where we know the full system
 - ▶ there are no hidden channels outside our algorithm
 - ▶ happens-before defined according to actual communication
 - ▶ concurrent events are as if spacelike
 - ▶ comparing their times can be meaningless

Logical clock condition

- ▶ Lamport defined a condition for any logical clock C


$$\text{if } a \rightarrow b \text{ then } C(a) < C(b)$$

- ▶ This is analogous to what we expect from physical clocks
 - ▶ any observer will compare timelike events in the same way
- ▶ Initially Lamport defined a clock value as a number
- ▶ It can be generalized to clock values in partially ordered sets

Lamport clocks

- ▶ A clock value is an integer, increasing along each causal path
- ▶ Each process i keeps a clock L_i
- ▶ An update (state transition) or send increments the clock
- ▶ A send attaches it to message m as L_m
- ▶ A receive of m stores one plus the maximum of clock values

$$L_i := \max(L_i, L_m) + 1$$



Advantages and limitations of Lamport clocks

- ▶ They have several advantages
 - ▶ cheap to compute, store and send in messages
 - ▶ can be used to define a total order of events

Defining a total order respecting happens-before

- ▶ assuming a total order of process identifiers
- ▶ for events a and b of processes i and j respectively

$$a < b \Leftrightarrow L(a) < L(b) \vee (L(a) = L(b) \wedge i < j)$$

- ▶ Lamport clocks only respect, not characterize happens-before
 - ▶ comparing clocks doesn't inform whether events are concurrent
 - ▶ except when the clocks happen to have exactly the same value

Characterizing happens-before

- ▶ Many different logical clocks can respect the clock condition
- ▶ For any logical clock

$$a \rightarrow b \Rightarrow C(a) < C(b)$$

- ▶ But the reverse is not true in general

$$C(a) < C(b) \not\Rightarrow a \rightarrow b$$

- ▶ A clock characterizes happens-before when the reverse is true

$$a \rightarrow b \Leftrightarrow C(a) < C(b)$$

- ▶ Such clocks need to have values in a partially ordered set


Causal history

- ▶ The causal history of an event is the set of all events that happened before and the event itself.

$$\mathbb{C}(e) = \{e\} \cup \{x \mid x \rightarrow e\}$$

- ▶ By definition

$$e \rightarrow e' \Leftrightarrow \mathbb{C}(e) \subset \mathbb{C}(e')$$



- ▶ Denoting i_n the n -th event of process i
- ▶ In the run, $\mathbb{C}(k_3) = \{i_1, i_2, j_1, j_2, j_3, j_4, j_5, k_2, k_3\}$

Vector clocks

- ▶ Causal histories are not to be used in practice
- ▶ But serve to reason about actual mechanisms (clocks)
- ▶ Looking at the previous example, to encode \mathbb{C}
 - ▶ if we number events sequentially in each process
 - ▶ it is enough to store the last event number from each process

$$V(k_3) = \{i \mapsto 2, j \mapsto 5, k \mapsto 3\}$$

- ▶ or if we number processes, simply a vector $[2, 5, 3]$

A *vector clock* is a vector (if processes are numbered) or map (in general) that maps process identifiers to their last event number present in a causal history.

Operations with vector clocks

- All events (update, send, receive) increment self component


$$V_i[i] := V_i[i] + 1$$

- A send attaches the vector clock to message m as V_m
- A receive also performs the pointwise maximum:

$$V_i := \max(V_i, V_m)$$


- Clocks compared by standard comparison of functions

$$V_i \leq V_j \Leftrightarrow \forall p \cdot V_i[p] \leq V_j[p]$$



Visual analogies for causal histories and logical clocks

- ▶ Over time many different logical clocks were designed
- ▶ Resorting to causal histories helps reasoning
- ▶ Clocks are different ways of representing causal histories
- ▶ A visual analogy helps
- ▶ For the same run with vector clocks



Variable number of nodes

- ▶ A vector clock using a map adapts to variable number of nodes
- ▶ We need only use unique node identifiers
- ▶ Each id not mapped is implicitly mapped to 0
- ▶ But this makes vector clocks acquire more entries over time
- ▶ Node retirement is tricky

Dynamic systems with process creation and retirement

- ▶ One possible execution model allows
 - ▶ forking a new node from an existing one
 - ▶ joining two nodes into one, retiring one node
 - ▶ state transitions
 - ▶ sending messages
- ▶ A clock for such dynamic systems could/should
 - ▶ have dynamic (plastic) representations of identities
 - ▶ a way to manage (split, gather, retire) identities
 - ▶ have the notion of the identity owned by a node

Interval Tree Clocks – a logical clock for dynamic systems

- ▶ A clock is a function defined over a tree of intervals
- ▶ Reasonably complex, easy to understand visually
 - ▶ light-gray area represents causal history
 - ▶ dark-gray area represents owned identity
 - ▶ updates inflate some light-gray area over the dark-gray
 - ▶ forks split dark-gray area, keep light-gray area
 - ▶ joins merge both areas
 - ▶ partial order is according to light-gray area inclusion
 - ▶ buddy areas can be merged to simplify representation

