

# State-based CRDTs

Paulo Sérgio Almeida



Universidade do Minho

## State-based Replication

Lattices and Order

State-based CRDTs

# Conflict-free Replicated Data Types

- ▶ Provide operations, like standard abstract datatypes
- ▶ Each datatype object replicated and accessed locally
  - ▶ Mutator operations update state
  - ▶ Query operations look at state and return result
- ▶ Information propagated to other replicas asynchronously
- ▶ Object highly available even under partitions


# The C from CRDTs

- ▶ C... Replicated Data Types
  - ▶ Convergent?
  - ▶ Conflict-free?
  - ▶ Commutative?
- ▶ Convergence while resolving conflicts
- ▶ Replicas keep converging; world does not have to stop
- ▶ Conflicts are dealt with semantically: spec of datatype
- ▶ Availability is achieved by forgoing total orders
- ▶ Concurrent operations will become visible in different orders
- ▶ Some confusion about what is commutative (spec vs impl)
  - ▶ some operations are not (semantically) commutative
  - ▶ effects of executing concurrent operations must be

# Operation-based vs state-based approaches


- ▶ State-based approaches
  - ▶ propagate replica states
  - ▶ detect mutual inconsistency
  - ▶ reconcile (merge) concurrent replicas
  - ▶ anti-entropy by opportunistic, “background” communication
  - ▶ can be made more incremental by delta-state approach
- ▶ Operation-based approaches
  - ▶ propagate information about operations
  - ▶ use a reliable messaging algorithm for propagation
  - ▶ need ordering guarantees (typically causal)
- ▶ Here we address state-based CRDTs

# State-based replication through state propagation



- ▶ Messages carry state, immediately delivered
- ▶ In figure, marks signal state change, labelled by resulting state
- ▶ States  $A \sqsubseteq B \sqsubseteq C \sqsubseteq D \sqsubseteq E$  evolve in sequence (even through different replicas)
- ▶ States  $E$  and  $G$  evolved concurrently, are conflicting ...  
... and must be reconciled (merged) to obtain  $H$
- ▶  $B = A$ ; also  $A \sqsubseteq C$ , so  $D = C$ ; also  $C \sqsubseteq I$ , so  $J = I$
- ▶ No state change when delivering  $n$  ( $B = A$ ),  $p$  (duplicate message) and  $q$  ( $F \sqsubseteq G \sqsubseteq H \sqsubseteq I = J$ )

# Detecting mutual inconsistency through version vectors



- ▶ Only updates increase self entry (not sends or receives)
- ▶ Send attaches *version vector* to message  $m$  as  $V_m$
- ▶ Receive performs pointwise maximum  $V_i := \max(V_i, V_m)$ 
  - ▶ assumes deterministic conflict resolution
  - ▶ otherwise, must also increase self entry
- ▶ VVs compared by standard comparison of functions

$$V_i \leq V_j \Leftrightarrow \forall p \cdot V_i[p] \leq V_j[p]$$

- ▶ Unlike with VCs, several nodes can have the same VV

# How to resolve conflicts

- ▶ Version vectors know whether two replicas conflict
- ▶ If that happens, some conflict resolution must be performed
- ▶ Classically performed in some ad hoc way
- ▶ Conflict resolution (merge) should be
  - ▶ deterministic: result as a function of inputs
  - ▶ obviously a commutative function
  - ▶ associative: same result when merging in different orders
  - ▶ monotonic: merging with “newer” state produces “newer” state

Mathematical concept of *join semilattices* is the solution



State-based Replication

Lattices and Order

State-based CRDTs

# Sets and partially ordered sets (posets)

- ▶ A set assumes no order between elements
- ▶ A partially ordered set (poset), is a set equipped with a binary relation  $\sqsubseteq$  which is:
  - ▶ (reflexive)  $p \sqsubseteq p$
  - ▶ (transitive)  $o \sqsubseteq p \wedge p \sqsubseteq q \Rightarrow o \sqsubseteq q$
  - ▶ (anti-symmetric)  $p \sqsubseteq q \wedge q \sqsubseteq p \Rightarrow p = q$
- ▶ Two unordered elements are called concurrent.
  - ▶ (concurrent)  $p \parallel q \Leftrightarrow \neg(p \sqsubseteq q \vee q \sqsubseteq p)$
- ▶ Some posets have *bottom*, an element smaller than any other:

$$\forall p \in P. \perp \sqsubseteq p$$

# Join semilattices

- ▶ In poset  $P$ , an upper bound of subset  $S$  is some  $u$  such that

$$\forall s \in S \cdot s \sqsubseteq u$$

- ▶ The *least upper bound* exists, when there is an upper bound smaller than any other upper bound
- ▶ In that case it is called the *join* of  $S$ , denoted  $\sqcup S$ .
- ▶ For two elements,  $\sqcup\{p, q\}$  is denoted by  $p \sqcup q$
- ▶ A poset  $P$  is a join semilattice if there exists a least upper bound for any pair of elements  $p$  and  $q$  in  $P$
- ▶ Properties of the join operator:
  - ▶ (idempotent)  $p \sqcup p = p$
  - ▶ (commutative)  $p \sqcup q = q \sqcup p$
  - ▶ (associative)  $o \sqcup (p \sqcup q) = (o \sqcup p) \sqcup q$

# Examples and non-examples of join semilattices

- ▶ Some join semilattices:
  - ▶ natural numbers  $\mathbb{N}$ ;  $\sqsubseteq = \leq$ ;  $\sqcup = \max$ ;  $\perp = 0$
  - ▶ booleans  $\mathbb{B}$ , with  $\text{False} \sqsubseteq \text{True}$ ;  $\sqcup = \vee$ ;  $\perp = \text{False}$
  - ▶ any totally ordered set (*chain*)
- ▶ Some posets which are not join semilattices:
  - ▶ unordered posets (*antichain*)
    - ▶ all elements being incomparable
  - ▶ strings under prefix ordering (e.g.,  $\text{small} \sqsubseteq \text{smallest} \parallel \text{smaller}$ )
    - ▶ all concurrent elements are not joinable

# Lattice compositions (1)

- ▶ Cartesian product  $A \times B$  of two join semilattices  $A$  and  $B$

$$\begin{aligned}(a_1, b_1) \sqsubseteq (a_2, b_2) &\Leftrightarrow a_1 \sqsubseteq a_2 \wedge b_1 \sqsubseteq b_2 \\ (a_1, b_1) \sqcup (a_2, b_2) &= (a_1 \sqcup a_2, b_1 \sqcup b_2)\end{aligned}$$

- ▶ Lexicographic product  $A \boxtimes B$  of two join semilattices  $A$  and  $B$ , when  $B$  has a bottom or  $A$  is a chain

$$\begin{aligned}(a_1, b_1) \sqsubseteq (a_2, b_2) &= a_1 \sqsubset a_2 \vee (a_1 = a_2 \wedge b_1 \sqsubseteq b_2) \\ (a_1, b_1) \sqcup (a_2, b_2) &= \begin{cases} (a_1, b_1) & \text{if } a_2 \sqsubset a_1 \\ (a_2, b_2) & \text{if } a_1 \sqsubset a_2 \\ (a_1, b_1 \sqcup b_2) & \text{if } a_1 = a_2 \\ (a_1 \sqcup a_2, \perp) & \text{if } a_1 \parallel a_2 \end{cases}\end{aligned}$$

## Lattice compositions (2)

- Powerset  $\mathcal{P}(S)$  for any set  $S$

$$\sqsubseteq = \subseteq \quad \sqcup = \cup$$

- Function space  $A \rightarrow B$  from set  $A$  to join semilattice  $B$

$$f \sqsubseteq g = \forall x \in A \cdot f(x) \sqsubseteq g(x) \quad (f \sqcup g)(x) = f(x) \sqcup g(x)$$

- Maps (partial functions)  $K \rightarrow V$ , to join semilattice  $V$  with bottom, where missing keys implicitly yield bottom

$$m(k) = \begin{cases} v & \text{if } (k, v) \in m \\ \perp & \text{otherwise} \end{cases}$$

are efficient representations of total function; same  $\sqsubseteq$  and  $\sqcup$

State-based Replication

Lattices and Order

State-based CRDTs

# State-based CRDTs


- ▶ Operations applied locally
- ▶ State is a join-semilattice
- ▶ Replicas send the full state
- ▶ Replicas merge received state using join operator  $\sqcup$
- ▶ Assume unreliable networks (loss, duplication, reordering)
- ▶ Sending full state ensures causal consistency



# Update functions and state mutators

- ▶ Update functions invoke a state mutator
- ▶ State mutator must be an *inflation*
  - ▶ (inflation)  $x \sqsubseteq f(x)$
  - ▶ (strict inflation)  $x \sqsubset f(x)$
- ▶ What is needed is that state evolves monotonically
  - ▶ when updates are locally invoked
  - ▶ when remote state is received and merged through join

# Inflations vs monotonic functions



- ▶ Erroneously, many places say mutators need to be monotonic
- ▶ State mutators need to be inflations
- ▶ Being monotonic is not necessary nor sufficient
- ▶  $f$  is an inflation but not monotonic
- ▶  $g$  is monotonic but not an inflation

## Grow only set $\text{GSet}\langle E \rangle$

$$\text{GSet}\langle E \rangle = \mathcal{P}(E)$$

$$\perp = \{\}$$

$$\text{add}_i(e, s) = s \cup \{e\}$$

$$\text{elements}(s) = s$$

$$\text{contains}(e, s) = e \in s$$

$$s \sqcup s' = s \cup s'$$

- ▶ Trivial CRDT: state same as sequential datatype
- ▶ Add already an inflation; mutator = update
- ▶ Merging replica states by set union
- ▶ Anonymous CRDT: no need for node ids in the state

# Single-writer principle and named CRDTs

- ▶ Powerful strategy: single writer principle
  - ▶ state partitioned in several parts
  - ▶ each node updates a part dedicated exclusively to itself
  - ▶ state is joined by joining respective parts
- ▶ Unique node ids can be used to partition state
  - ▶ as a map from ids to parts
  - ▶ example: version vectors
- ▶ CRDTs that use node ids in the state are *named CRDTs*

# State-based PCounter

$$\text{GCounter} = \mathbb{I} \rightarrow \mathbb{N}$$

$$\perp = \{\}$$

$$\text{inc}_i(m) = m\{i \mapsto m(i) + 1\}$$

$$\text{value}(m) = \sum_{j \in \mathbb{I}} m(j)$$

$$m \sqcup m' = \{j \mapsto \max(m(j), m'(j)) \mid j \in \text{dom } m \cup \text{dom } m'\}$$

- ▶ State maps replica ids to integers
- ▶ inc increments self entry ( $i$ )
- ▶ Merge is pointwise max
- ▶ Similar in structure to a version-vector

# State-based PNCounter

$$\begin{aligned}\text{PNCounter} &= \text{GCounter} \times \text{GCounter} \\ \perp &= (\perp, \perp) \\ \text{inc}_i((p, n)) &= (\text{inc}_i(p), n) \\ \text{dec}_i((p, n)) &= (p, \text{inc}_i(n)) \\ \text{value}((p, n)) &= \text{value}(p) - \text{value}(n) \\ (p, n) \sqcup (p', n') &= (p \sqcup p', n \sqcup n')\end{aligned}$$

- ▶ Problem: dec is not natively an inflation
- ▶ Solved through a pair of GCounters (product composition)
  - ▶ increments and decrements tracked separately
  - ▶ counter value obtained as difference
- ▶ In practice, a single map to pairs is used

# How to forget without tombstones

- ▶ Many times we want to remove things
  - ▶ example: set with add and remove
- ▶ But must always use inflations
  - ▶ cannot use single set for state and remove elements
- ▶ Using tombstones to mark removal makes state grow forever
- ▶ Example: two-phase set
  - ▶ allows add and remove
  - ▶ once removed, an element cannot be re-added
  - ▶ implementation with a pair of sets, to track adds and removes
  - ▶ state is always growing
- ▶ Tombstones can be avoided using *Causal CRDTs*


# Causal CRDTs

- ▶ State has two components:
  - ▶ a *dot store*
  - ▶ a *causal context*
- ▶ The dot store (DS)
  - ▶ is a container for datatype-specific information
  - ▶ information tagged with unique event identifiers: *dots*
- ▶ The causal context (CC)
  - ▶ represents causal history: ids of all visible updates
  - ▶ in state-based CRDTs encoded by a version vector

Any information covered by the causal context and not present in the dot store has already been removed



# Joining causal CRDTs



- ▶ The dot store that results from a join of  $x$  and  $y$ 
  - ▶ has dots from  $x$ 's DS not covered by  $y$ 's CC
  - ▶ has dots from  $y$ 's DS not covered by  $x$ 's CC
  - ▶ has the dots present in both DSs
- ▶ The causal context upon a join is the join of causal contexts

## Some dot stores

DS

$$\text{DotSet} : \text{DS} = \mathcal{P}(\mathbb{I} \times \mathbb{I})$$

$$\text{DotFun}\langle V : \text{Lattice} \rangle : \text{DS} = \mathbb{I} \times \mathbb{I} \rightarrow V$$

$$\text{DotMap}\langle K, V : \text{DS} \rangle : \text{DS} = K \rightarrow V$$

- ▶ DotMap in particular is a powerful concept
- ▶ Allows building map CRDTs which embed CRDTs, or maps which embed maps ... which embed CRDTs

$$\text{DotMap}\langle K_1, \text{DotMap}\langle K_2, \text{DotFun}\langle \mathcal{P}(E) \rangle \rangle \rangle$$

# Causal CRDTs

$$\text{Causal}\langle T : \text{DS} \rangle = T \times \text{CausalContext}$$

$$\sqcup : \text{Causal}\langle T \rangle \times \text{Causal}\langle T \rangle \rightarrow \text{Causal}\langle T \rangle$$

when  $T : \text{DotSet}$

$$(s, c) \sqcup (s', c') = ((s \cap s') \cup (s \setminus c') \cup (s' \setminus c), c \cup c')$$

when  $T : \text{DotFun}\langle \_ \rangle$

$$(m, c) \sqcup (m', c') = (\{k \mapsto m(k) \sqcup m'(k) \mid k \in \text{dom } m \cap \text{dom } m'\} \cup \\ \{(d, v) \in m \mid d \notin c'\} \cup \{(d, v) \in m' \mid d \notin c\}, c \cup c')$$

when  $T : \text{DotMap}\langle \_, \_ \rangle$

$$(m, c) \sqcup (m', c') = (\{k \mapsto v(k) \mid k \in \text{dom } m \cup \text{dom } m' \wedge v(k) \neq \perp\}, c \cup c') \\ \text{where } v(k) = \text{fst}((m(k), c) \sqcup (m'(k), c'))$$

# Multi-value register MVReg $\langle E \rangle$

$$\begin{aligned}\text{MVReg}\langle E \rangle &= \text{Causal}\langle \text{DotFun}\langle E_{\perp}^{\top} \rangle \rangle \\ \text{write}_i(v, (m, c)) &= (\{(i, c[i] + 1) \mapsto v\}, c\{i \mapsto c[i] + 1\}) \\ \text{read}((m, c)) &= \{v \mid (k, v) \in m\}\end{aligned}$$

- ▶ Dot store is a DotFun: map from dots to values in  $E$
- ▶ Value is range of DotFun
- ▶ Previously defined join for the DotFun case used
  - ▶ value mapped from a given dot remains immutable
  - ▶ therefore, no join of values for any given key (dot) happens
- ▶ Join keeps concurrently written values

## Observed-remove set $\text{ORSet}\langle E \rangle$

$$\text{ORSet}\langle E \rangle = \text{Causal}\langle \text{DotMap}\langle E, \text{DotSet} \rangle \rangle$$

$$\text{add}_i(e, (m, c)) = (m\{e \mapsto \{(i, c[i] + 1)\}\}, c\{i \mapsto c[i] + 1\})$$

$$\text{remove}_i(e, (m, c)) = (\{e\} \triangleleft m, c)$$

$$\text{elements}((m, c)) = \text{dom } m$$

- ▶ Add replaces, for that key, any set of dots by a new singleton
- ▶ Remove simply removes that key (domain subtraction)
  - ▶ no new event introduced in the causal context
- ▶ Previously defined join for the DotMap case used
  - ▶ keeps concurrently added elements
  - ▶ removes elements removed elsewhere if no dot survives